

Analytics in a Day

Azure Synapse + Power BI better together

Exercise 07

Optimize a Power BI Model

Overview

The estimated time to complete this exercise is 15 minutes.

*Important: It's a prerequisite that you successfully complete **Exercise 06** before commencing this exercise.*

In this exercise, you are working in the role of a **data architect** or **BI developer**.

You will use Power BI Desktop to continue the development of the data model you developed in **Exercise 06**. The goal in this exercise is to improve query performance. You will commence by configuring dimension table storage as dual. You will then add an aggregation table to accelerate query performance. The model storage will be switched to mixed mode: Some tables will cache data to further boost query performance. You will finalize the exercise by deploying the model to the Power BI service, so it is ready for reporting in **Exercise 08**.


Section 1: Add an Aggregation Table

In this section, you will configure dimension tables as dual storage mode table. You will then create an aggregation table to boost Power BI query performance for date, geography, and profit reporting.

Task 1: Configure Dual Storage

In this task, you will configure dual storage for all dimension tables.

1. Switch to the Power BI Desktop solution you developed in **Exercise 06**.
2. Switch to Model view.
3. While pressing the **Ctrl** key, multi-select each of the five dimension tables:
 - Customer
 - Date
 - Geography
 - Product
 - Salesperson
4. In the **Properties** pane, from within the **Advanced** section, in the **Storage Mode** dropdown list, select **Dual**.



It's common to set dimension tables to use dual storage mode. This way, when used by report slicers, they deliver fast performance.

5. When prompted to set the storage mode, click **OK**.




- When the refresh completes, save the Power BI Desktop solution.

The data model is now in mixed mode. It's a composite model consisting of DirectQuery storage mode tables and import storage mode tables.

- In File Explorer, notice the file size has grown as a result of the imported data for the dimension tables.

When the model stores data, you need to ensure the cached data is current. The model must be refreshed on a frequent basis to ensure import data is in sync with the source data.

- In Report view, in the status bar, at the bottom-right, notice that the storage mode is now mixed.

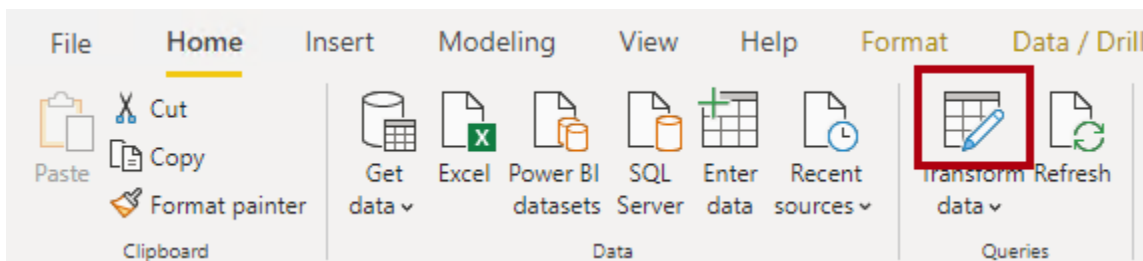


- In the **Performance Analyzer** pane, start recording, and then refresh visuals.
- Notice that the query result for the slicer is now sub-second.
- In the **Performance Analyzer** pane, stop recording.

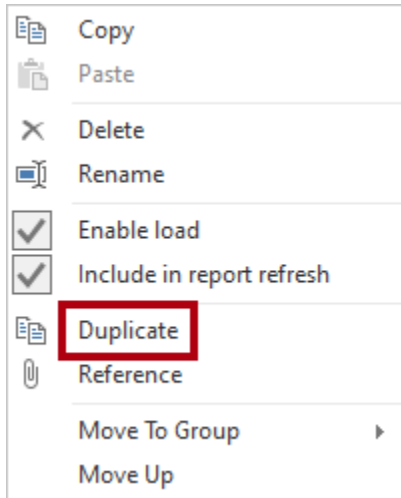
Task 2: Create an Aggregation Table

In this task, you will create an aggregation table to accelerate Power BI report visuals that specifically query by date and geography, and summarize profit.

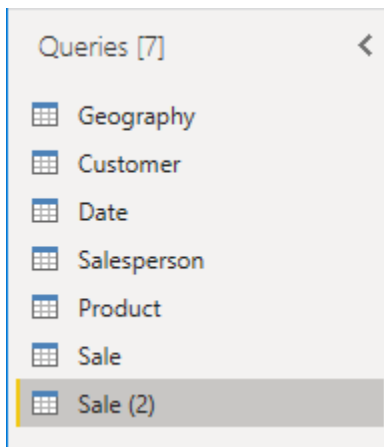
- To open the Power Query Editor window, on the **Home** ribbon tab, from inside the **Queries** group, click the **Transform Data** icon.



2. In the Power Query Editor window, from inside the **Queries** pane, right-click the **Sale** query, and then select **Duplicate**.




3. In the **Queries** pane, notice the addition of a new query.




You'll apply a transformation to group by the **CityKey** and **InvoiceDateKey** columns, and aggregate the sum of **Profit Amount** column.

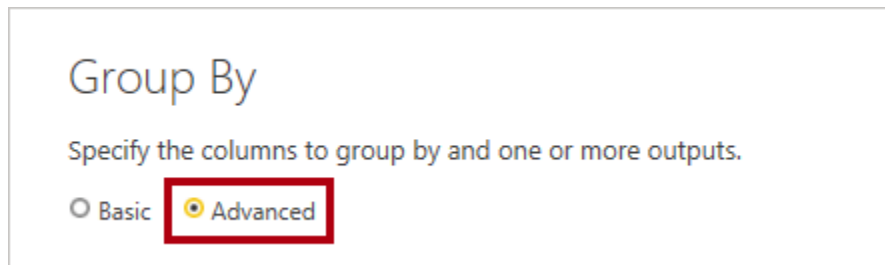
4. Rename the query as **Sale Agg**.



5. On the **Transform** ribbon tab, from inside the **Table** group, click **Group By**.

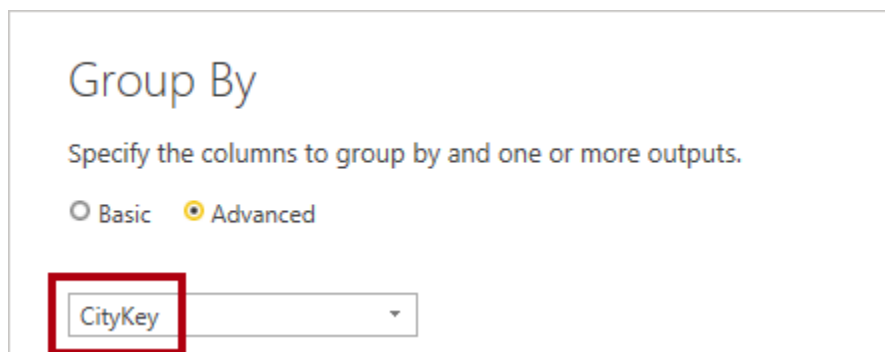


6. In the **Group By** window, select the **Advanced** option.




The advanced option allows grouping by more than one column.

7. In the grouping dropdown list, ensure that **CityKey** is selected.




8. Click **Add Grouping**.
9. In the second grouping dropdown list, select **InvoiceDateKey**.
10. In the **New Column Name** box, replace the text with **Profit Amount**.
11. In the **Operation** dropdown list, select **Sum**.
12. In the **Column** dropdown list, select **Profit Amount**.



13. Click **OK**.



14. On the **Home** ribbon tab, from inside the **Close** group, click the **Close & Apply** icon.



A new table is added to the model.

15. Save the Power BI Desktop solution.

Task 3: Configure Aggregations

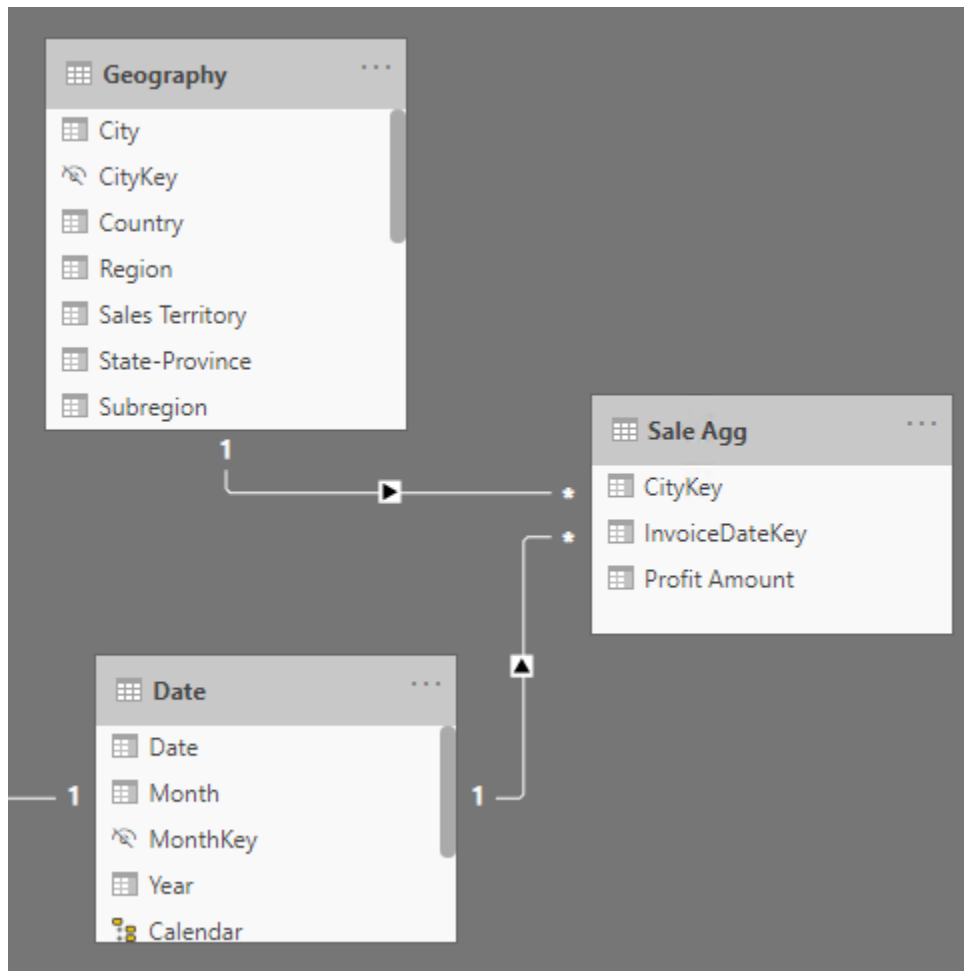
In this task, you will switch the aggregation table to import data. You will then create model relationships to the aggregation table and manage aggregations.

1. Switch to Model view.
2. Position the **Sale Agg** table so that it is near the **Geography** and **Date** tables.
3. Set the storage mode for the **Sale Agg** table as **Import**.

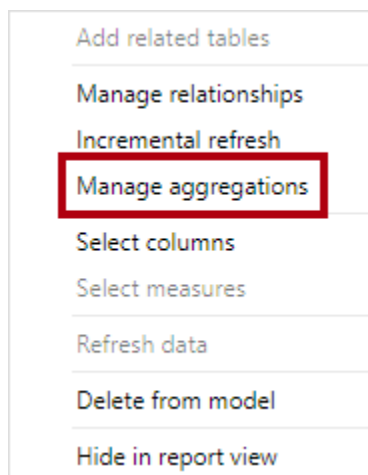


4. Create two model relationships:

- Relate the **Sale Agg** table **CityKey** column to the **Geography** table **CityKey** column
- Relate the **Sale Agg** table **InvoiceDateKey** column to the **Date** table **Date** column



5. Right-click the **Sale Agg** table, and then select **Manage Aggregations**.



6. In the **Manage Aggregations** window, for the **Profit Amount** aggregation column, set the following properties:
 - Summarization: **Sum**
 - Detail table: **Sale**
 - Detail column: **Profit Amount**

AGGREGATION COLUMN	SUMMARIZATION	DETAIL TABLE	DETAIL COLUMN
CityKey	Select Summarizatio...		
InvoiceDateKey	Select Summarizatio...		
Profit Amount	Sum	Sale	Profit Amount

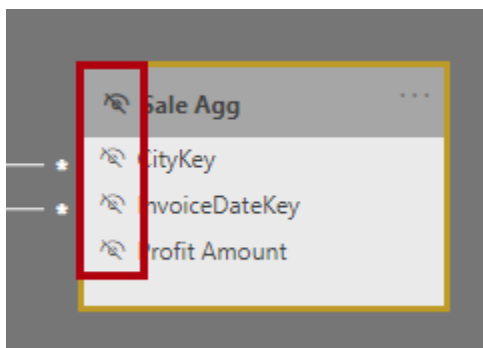
This table will be hidden if aggregations are set because aggregation tables must be hidden.

7. Notice the warning that describes the table will be hidden.

*The table will be hidden in a different way to other hidden model objects (like the key columns you hid in **Exercise 06**). Aggregation tables are always hidden, and they can't even be referenced in model calculations.*

8. Click **Apply All**.

9. In the model diagram, notice that the **Sale Agg** table is now hidden.



10. In the model diagram, select the **Sale Agg** table.
11. Switch to Report view.

12. In the **Performance Analyzer** pane, start recording, and then refresh visuals.

13. Notice that the query results for the table visual is now sub-second.

*Because the **Geography** and **Date** tables use dual storage mode, when a report visual queries them at the same time as the aggregation table, Power BI will query the model cache. There's no need to use DirectQuery to query the data.*

14. In the **Performance Analyzer** pane, stop recording.


Section 2: Publish the Model

In this section, you will publish the model and complete some post-publication tasks.

Task 1: Publish the Model

In this task, you will publish the model.

1. In Power BI Desktop, on the **Home** ribbon tab, from inside the **Share** group, click **Publish**.



2. If prompted to save changes, click **Yes**.
3. In the **Publish to Power BI** window, select the lab workspace (do not use **My Workspace**).
4. Click **Select**.



5. When publication has completed, click **Got It**.
6. Close Power BI Desktop.

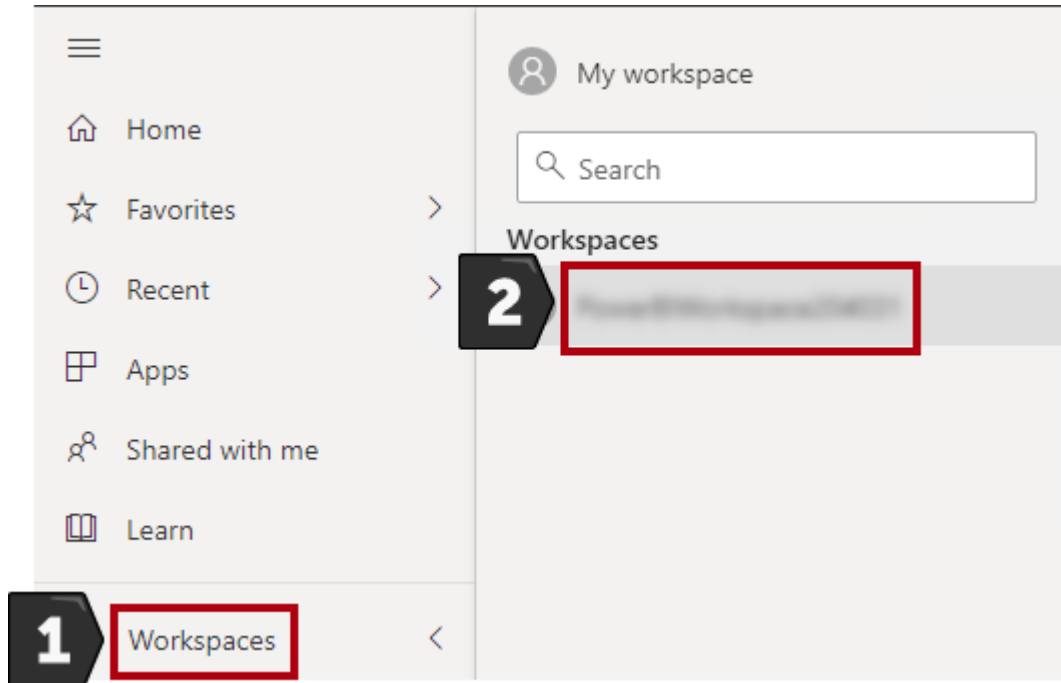
*You will open a new instance of Power BI Desktop in **Exercise 08** when you create a report.*

Task 2: Complete Post-Publication Tasks


In this task, you will complete some post-publication tasks.

You'll complete the post-publication tasks using the Power BI service because it's not possible to do them in Synapse Studio.

1. In the Power BI web browser session, open your lab workspace.



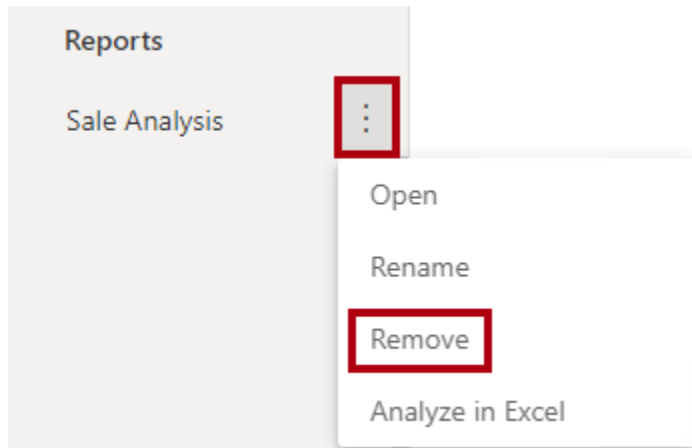
2. In the **Navigation** pane (located at the left), expand the workspace, and then verify that the **Sale Analysis** dataset exists.



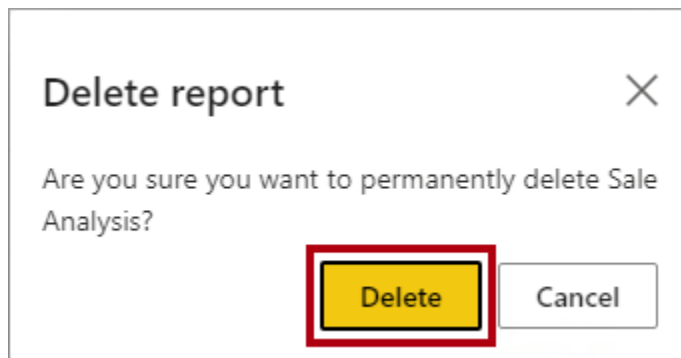
3. Notice there is also the **Sale Analysis** report.

*It wasn't our intention to publish a report, it was published alongside the model. You'll develop a report in **Exercise 08**. So, we'll delete this report.*


4. In the **Navigation** pane, hover the cursor over the **Sale Analysis** report, click the vertical ellipsis (...), and then select **Remove**.



5. When prompted to delete the report, click **Delete**.

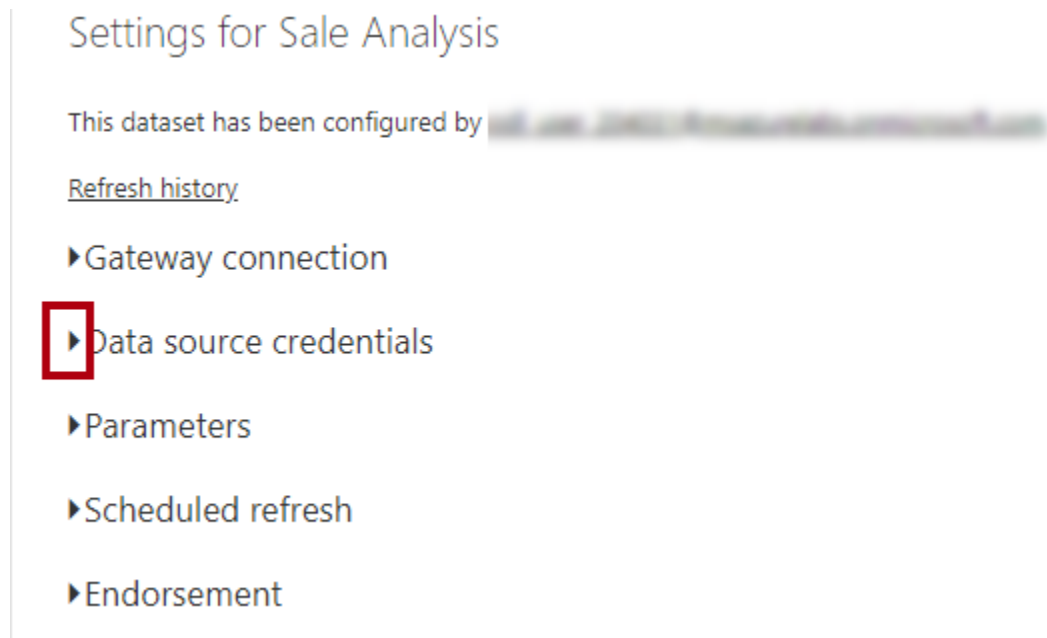


6. To apply data source credentials, in the **Navigation** pane, hover the cursor over the **Sale Analysis** dataset, click the vertical ellipsis, and then select **Settings**.



- Expand the **Data Source Credentials** section.

You'll see an error, and it's expected. You will address the error in the next step.



- To assign credentials, click the **Edit Credentials** link.
- In the window, in the **Authentication Method** dropdown list, ensure that **Microsoft Account** is selected.
- In the **Privacy Level** dropdown list, select **Organizational**.

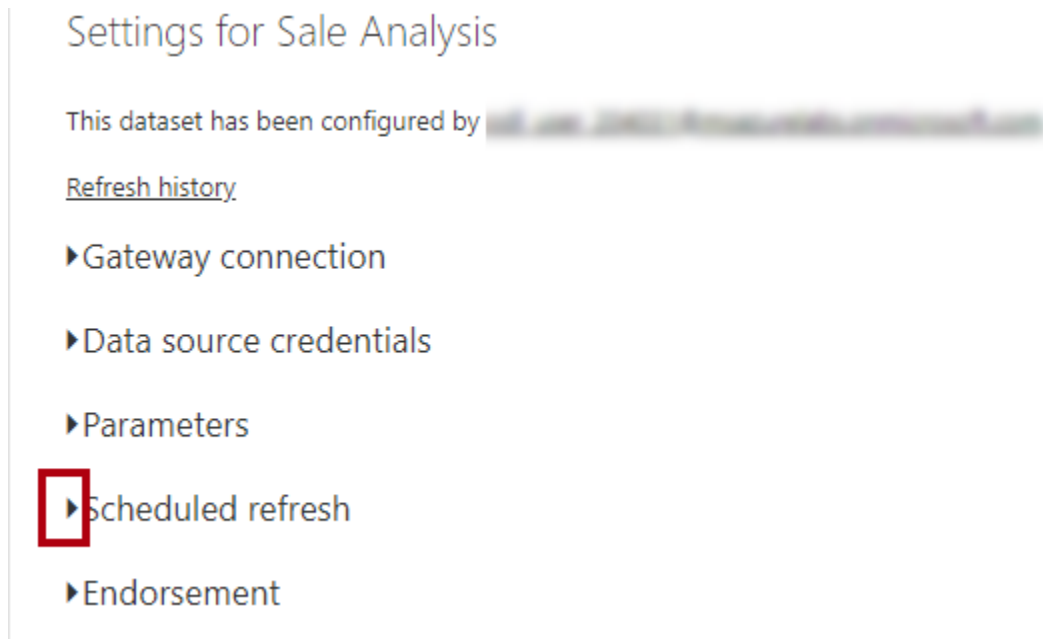
If you need the Power BI report user identity to flow to Azure Synapse (because per-user access permission must be enforced), you can check the checkbox. When the checkbox is left unchecked, the identity you will use to sign in (at the next step) will be used for all connections.

- Click **Sign In**.



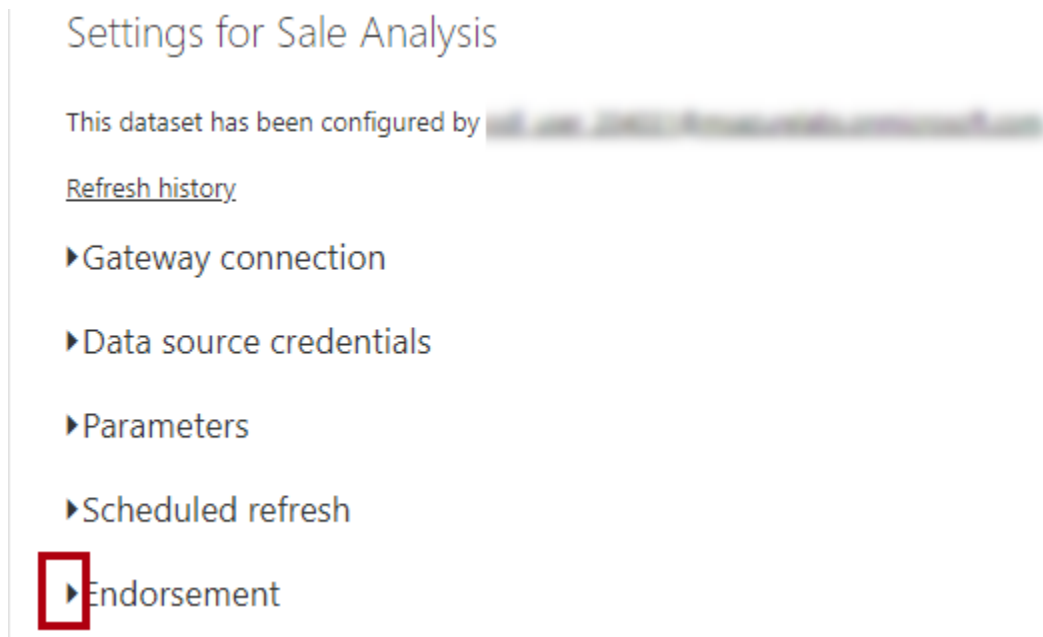
- Use the lab Azure credentials to sign in.

13. Expand the **Scheduled Refresh** section.



In this exercise, you won't schedule data refresh because it's a lab. In the real world, because your dataset contains import data (for the dimension tables and the aggregation table), you would schedule data refresh to keep the aggregation and dimension table import data current. It's possible, too, that your Azure Data Factory pipelines could send refresh commands using the Power BI REST API, once the data warehouse load has completed.

14. Expand the **Endorsement** section.



15. Select the **Promoted** option.

*The promoted endorsement communicates that the model is production-ready. In this exercise, ideally, you'd select the **Certified** option. A certified dataset is one that's truly reliable and authoritative, designed for use across the organization. (It's not available for your trial account.)*

16. In the **Description** box, enter: **Lab dataset**

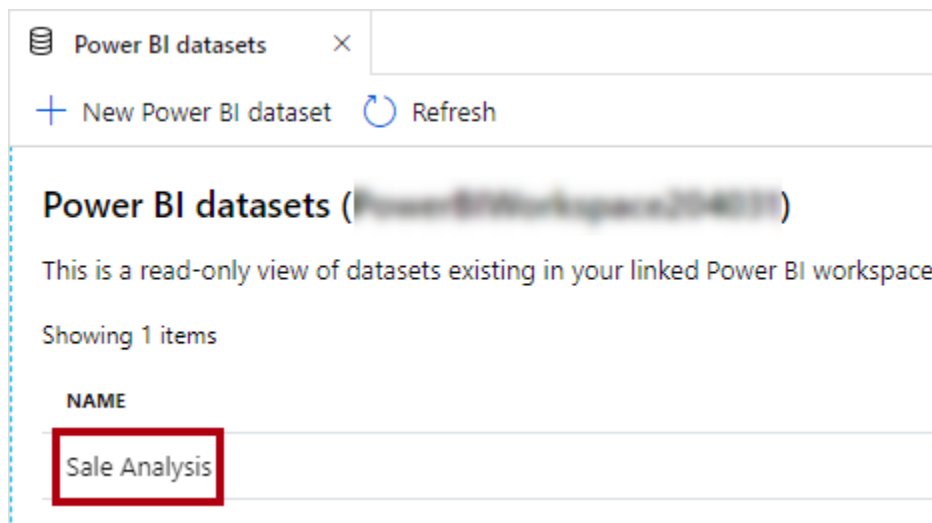
17. Click **Apply**.



18. Switch to the Azure Synapse web browser session.

19. In the **Develop** hub, select **Power BI Datasets**.

20. Notice that the **Sale Analysis** dataset is listed.



*The dataset is published and is configured ready for use. In **Exercise 08**, you'll perform a live connection to the dataset and create a report.*

Summary

In this exercise, you used Power BI Desktop to configure dimension table storage as dual. You then added an aggregation table to accelerate query performance. The model storage was switched to mixed mode: Some tables now cache data to further boost query performance. You finalized the exercise by deploying the model to the Power BI service, so it is ready for reporting in **Exercise 08**.

Terms of Use

© 2020 Microsoft. All rights reserved.

By using this hands-on lab, you agree to the following terms:

The technology/functionality described in this hands-on lab is provided by Microsoft Corporation in a “sandbox” testing environment for purposes of obtaining your feedback and to provide you with a learning experience. You may only use the hands-on lab to evaluate such technology features and functionality and provide feedback to Microsoft. You may not use it for any other purpose. Without written permission, you may not modify, copy, distribute, transmit, display, perform, reproduce, publish, license, create derivative works from, transfer, or sell this hands-on lab or any portion thereof.

COPYING OR REPRODUCTION OF THE HANDS-ON LAB (OR ANY PORTION OF IT) TO ANY OTHER SERVER OR LOCATION FOR FURTHER REPRODUCTION OR REDISTRIBUTION WITHOUT WRITTEN PERMISSION IS EXPRESSLY PROHIBITED.

THIS HANDS-ON LAB PROVIDES CERTAIN SOFTWARE TECHNOLOGY/PRODUCT FEATURES AND FUNCTIONALITY, INCLUDING POTENTIAL NEW FEATURES AND CONCEPTS, IN A SIMULATED ENVIRONMENT WITHOUT COMPLEX SET-UP OR INSTALLATION FOR THE PURPOSE DESCRIBED ABOVE. THE TECHNOLOGY/CONCEPTS REPRESENTED IN THIS HANDS-ON LAB MAY NOT REPRESENT FULL FEATURE FUNCTIONALITY AND MAY NOT WORK THE WAY A FINAL VERSION MAY WORK. WE ALSO MAY NOT RELEASE A FINAL VERSION OF SUCH FEATURES OR CONCEPTS. YOUR EXPERIENCE WITH USING SUCH FEATURES AND FUNCTIONALITY IN A PHYSICAL ENVIRONMENT MAY ALSO BE DIFFERENT.

FEEDBACK If you give feedback about the technology features, functionality and/or concepts described in this hands-on lab to Microsoft, you give to Microsoft, without charge, the right to use, share and commercialize your feedback in any way and for any purpose. You also give to third parties, without charge, any patent rights needed for their products, technologies and services to use or interface with any specific parts of a Microsoft software or service that includes the feedback. You will not give feedback that is subject to a license that requires Microsoft to license its software or documentation to third parties because we include your feedback in them. These rights survive this agreement.

MICROSOFT CORPORATION HEREBY DISCLAIMS ALL WARRANTIES AND CONDITIONS WITH REGARD TO THE HANDS-ON LAB, INCLUDING ALL WARRANTIES AND CONDITIONS OF MERCHANTABILITY, WHETHER EXPRESS, IMPLIED OR STATUTORY, FITNESS FOR A PARTICULAR PURPOSE, TITLE AND NON-INFRINGEMENT. MICROSOFT DOES NOT MAKE ANY ASSURANCES OR REPRESENTATIONS WITH REGARD TO THE ACCURACY OF THE RESULTS, OUTPUT THAT DERIVES FROM USE OF THE VIRTUAL LAB, OR SUITABILITY OF THE INFORMATION CONTAINED IN THE VIRTUAL LAB FOR ANY PURPOSE.

DISCLAIMER This lab contains only a portion of new features and enhancements in Microsoft Power BI. Some of the features might change in future releases of the product.